

Training loss over 30 epochs, 4 nodes / 16 GPUs

Rank-averaged mean diffusion loss per epoch. 960 steps, 1411 s.

No trend: the training subset predates OpenDDE's 2021-09 data cutoff.

